

Parallel Graph Coloring Using MPI

CSC 770 – Parallel Computing

Term Project

Presented by: Sanaria & Ben

Algorithm 1: Partitioning Strategy

Parallel greedy coloring with priority-based conflict resolution

Algorithm 2: Divide-and-Conquer Strategy

Maximal Independent Set (MIS)-based approach

</> Implemented in C with MPI



Graph Coloring Algorithms

Project Overview

Graph Coloring Algorithms

Parallel Computing - CSC 770

Project Details



Topic chosen from the course project list: [Graph Coloring](#)



Goal: **design two different parallel algorithms** for the problem



Algorithms in this project:

- Algorithm 1: [Partitioning Strategy](#)
- Algorithm 2: [Divide-and-Conquer Strategy](#)



Main focus of this work:

- Explain how each algorithm colors the graph in parallel
- Compare their performance as the number of processes increases

Graph Coloring Problem

Problem Definition

Given: A graph $G=(V,E)$

Assign: A color $C(v)$ to each vertex $v \in V$

Constraint: If $(u,v) \in E$, then $C(u) \neq C(v)$

Goal: Find a valid coloring

Prefer fewer colors, but optimal minimum is NP-hard

Applications:



Register allocation in compilers

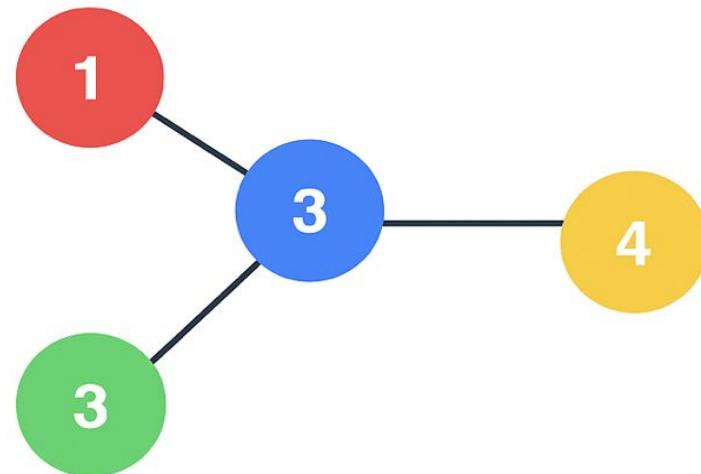


Timetable and exam scheduling



Frequency assignment in networks

Example



Valid coloring: Each vertex has a different color from its neighbors.

Uses 4 colors: Not optimal, but valid.

Why Parallel Graph Coloring?



Problem Context

Real-world graphs can be **very large**

Sequential coloring may be **too slow** for practical applications

Parallelism Benefits



Faster coloring for large graphs



Using multiple cores / processors

Theoretical Models



PRAM Model

Parallel Random Access Machine

EREW

CREW

CRCW

Practical Implementation



MPI Message Passing

Implementation model

Designing with PRAM-style thinking



Implementing with MPI in C

Background: PRAM & MPI



PRAM

Parallel Random Access Machine

- ✓ Many processors with shared memory
- ✓ Synchronous execution model
- ✓ Our two algorithms apply PRAM-style ideas but are implemented using MPI.

Submodels:

EREW

Exclusive Read Exclusive Write

CREW

Concurrent Read Exclusive Write

CRCW

Concurrent Read Concurrent Write



MPI

Message Passing Interface

- ✓ Each process has its own memory
- ✓ Processes have ranks in MPI_COMM_WORLD

Basic Functions:



MPI_Init



MPI_Finalize



MPI_Comm_size



MPI_Comm_rank



MPI_Send



MPI_Recv

Graph Representation & Data Structures

Graph Representation

- n = number of vertices
- $\text{adj}[i][j] = 1$ if edge between i and j , 0 otherwise

	v0	v1	v2	v3	v4
v0					

Arrays Used in Both Algorithms

 $\text{color}[v]$ – current color (0 = uncolored)

 $\text{priority}[v]$ – random weight for tie-breaking

MPI Partitioning

Vertices divided into blocks among processes:

Process r handles vertices from start_r to end_r

Extra for Divide-and-Conquer

 $\text{active}[v]$ – 1 if still in remaining graph, 0 if already colored

Two Parallel Graph Coloring Strategies



Algorithm 1: Partitioning Strategy



Vertices partitioned among processes

- ✓ Partition vertices among MPI processes
- ✓ Parallel greedy coloring with priorities
- ✓ Each process colors its assigned vertices
- ✓ Uses random priorities to break ties



Algorithm 2: Divide-and-Conquer Strategy



Maximal Independent Sets (MIS) coloring

- ✓ Repeatedly find Maximal Independent Sets (MIS)
- ✓ Color each MIS with a new color
- ✓ Remove colored vertices from graph
- ✓ Repeat until all vertices are colored

Both Algorithms Use:



C + MPI



Random priorities



Communication to sync colors

Algorithm 1: Partitioning Strategy – Intuition

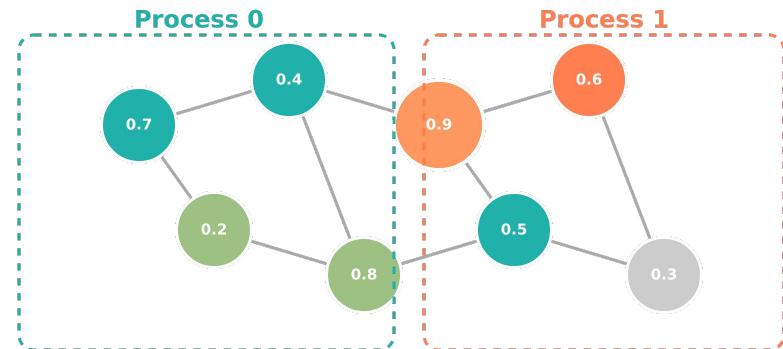
Algorithm Intuition

Each process manages a subset of vertices in the graph, allowing for parallel processing.

Key Steps:

-  **Local Processing:** Each process works on its assigned vertices
-  **Priority Check:** Verify if vertex has highest priority among neighbors
-  **Color Selection:** Choose smallest available color greedily
-  **Repeat:** Continue until all vertices are colored

Visual Intuition



 **Key Insight:** Priorities prevent conflicts while allowing parallel progress.

Algorithm 1: Partitioning Strategy – Steps

Input & Broadcast

Partition

Initialize

Repeat Until Done



Input & Broadcast

- Rank 0 reads n and adjacency matrix from file
- Uses MPI_Bcast to send graph to all processes



Partition Vertices

- Compute start and end indices for each process
- Each process handles a subset of vertices



Initialization

- For all vertices: $\text{color}[v] = 0$
- For all vertices: $\text{priority}[v] = \text{random}(0,1)$



Repeat Until Done

- For each local uncolored vertex: Check uncolored neighbors
- If priority is highest: choose smallest unused neighbor color and assign
- Synchronize $\text{color}[]$ across processes (e.g., MPI_Allgather) and use MPI_Reduce / MPI_Bcast

Algorithm 1: Communication & Time Complexity

Communication Patterns



MPI_Bcast

Broadcast graph size and adjacency matrix



MPI_Allgather

Share updated colors across all processes



MPI_Reduce + MPI_Bcast

Check if any vertex was colored this round

Time Complexity Analysis

Definitions:

- n = vertices
- Δ = max degree
- p = number of processes
- R = number of rounds

Local work per process:

$$\sim (n/p) \cdot \Delta \cdot R$$

Total time:

$$T_1(p) \approx O((n\Delta/p) \cdot R + R \cdot \text{comm_cost})$$

Algorithm 2: Divide-and-Conquer (MIS-Based) – Intuition

Maximal Independent Set (MIS)



Independent Set

No two vertices are adjacent (no edges between them)



Maximal

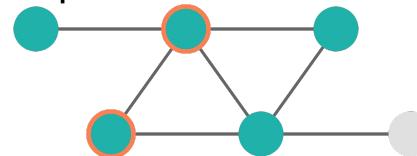
Cannot add more vertices without breaking independence

 In each round, we find an MIS in parallel, color its vertices, and remove them from the graph.

Algorithm Intuition

- 1 Find an MIS in parallel
- 2 Color all vertices in MIS with current color
- 3 Remove colored vertices ($\text{active}[v] = 0$)
- 4 Repeat on remaining graph with new color

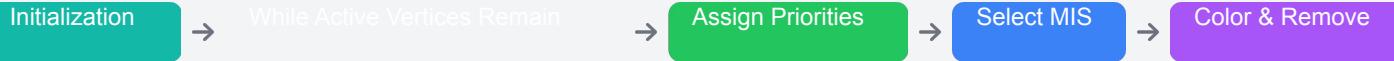
Original Graph



MIS (highlighted)

After Coloring and Removal

Algorithm 2: Divide-and-Conquer – Steps



1. Initialization

For all vertices:

- $\text{color}[v] = 0$ (uncolored)
 - $\text{active}[v] = 1$ (active)
- $\text{current_color} = 1$

2. While Active Vertices Remain

- Each process counts active vertices $\rightarrow \text{local_active}$
- $\text{MPI_Reduce} \rightarrow \text{global_active}$
- If $\text{global_active} == 0$, stop

3. Assign Priorities

- For active vertices: $\text{priority}[v] = \text{random}(0,1)$
- Sync priorities with MPI_Allgather

4. Select MIS in Parallel

- For each local active vertex v :
- If no active neighbor has higher priority $\rightarrow v$ joins MIS

5. Color MIS & Remove

- If v in MIS: $\text{color}[v] = \text{current_color}$
 - If v in MIS: $\text{active}[v] = 0$
 - Sync $\text{color}[]$ and $\text{active}[]$ with MPI_Allgather
- $\text{current_color}++$

Algorithm 2: Communication & Time Complexity

Communication Patterns



MPI_Bcast

Graph data at the beginning



MPI_Allgather

Synchronize: priority[], inMIS[], color[], active[]



MPI_Reduce

Sum of active vertices to detect termination

Time Complexity Analysis

Parallel Complexity (high-level):

$$T_2(p) \approx O((n\Delta/p) \cdot R' + R' \cdot \text{comm_cost})$$

- Let R' = number of color rounds (MIS selections)
- Each round: Work $\sim (n/p) \cdot \Delta$
- Plus communication for synchronization



Why Algorithm 2 is typically slower:

$R' > R \rightarrow$ more rounds than Algorithm 1 \rightarrow usually slower

Experimental Setup



Implementation

- ✓ Language: **C**
- ✓ Parallel library: **MPI**
- ✓ Platform: **Multi-core CPU**



Graphs Used

Random graphs with vertices:

n = 1000

n = 2000

n = 5000



Processes Tested

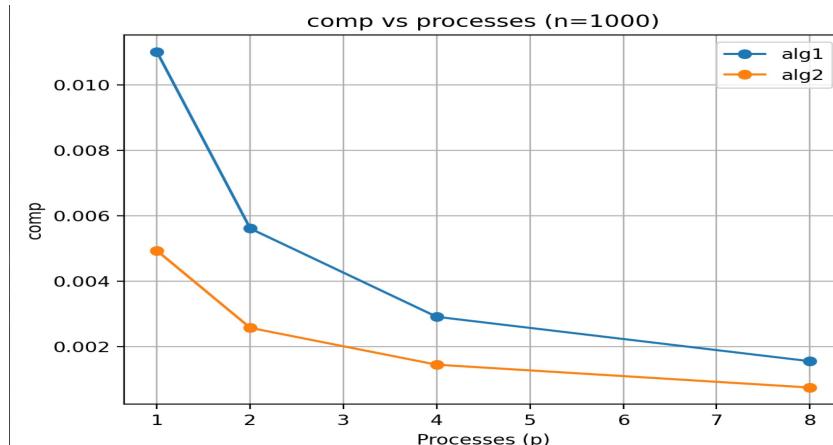
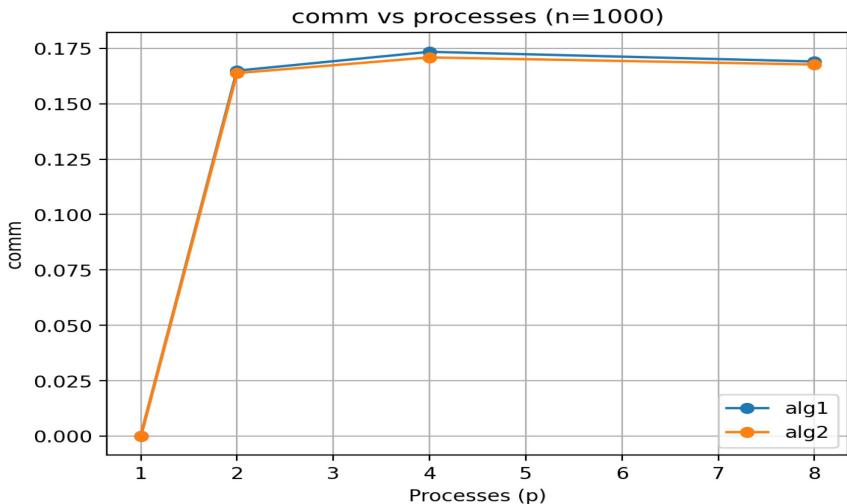
p = 1 **p = 2** **p = 4** **p = 8**



Metrics Collected

- ⌚ Execution time (seconds)
- ⌚ Speedup = $T(1)/T(p)$
- ⌚ Number of colors used

Results – Execution Time vs Processes



Key Observations

Algorithm Comparison

Partitioning Strategy shows consistently lower runtime with more processes

Communication Overhead

Overhead limits improvement at higher process counts

Performance Trend

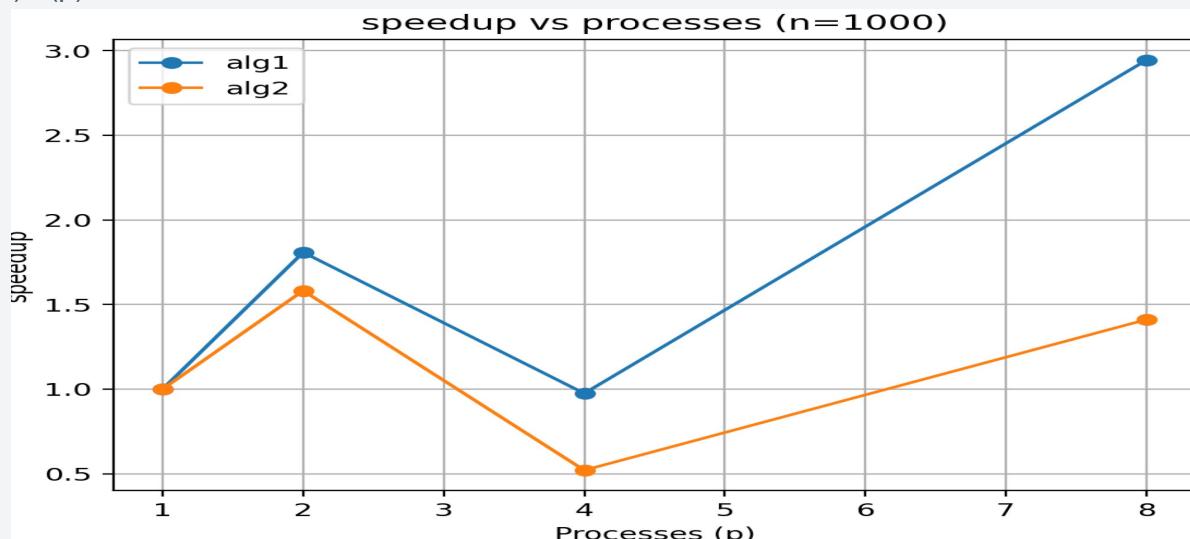
Execution time decreases as processes increase for both algorithms

Results – Speedup vs Processes



Speedup Definition:

$$\text{Speedup}(p) = T(1)/T(p)$$



Key Observations:

Nearly Linear Speedup

For small # of processes, both algorithms show linear speedup

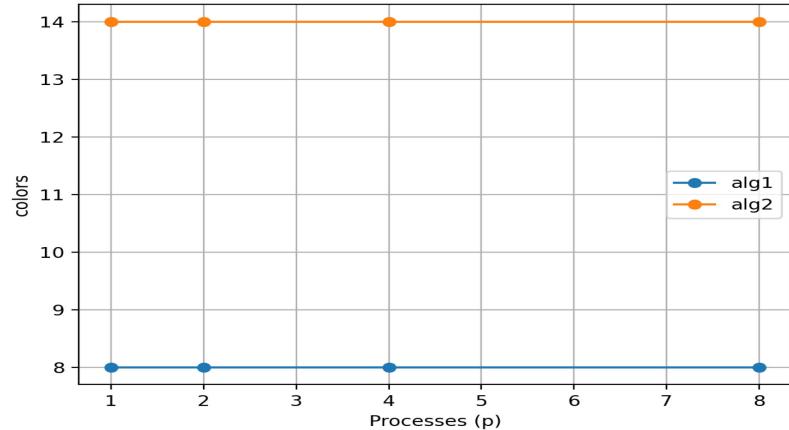
Communication Dominance

Speedup flattens as communication overhead becomes significant

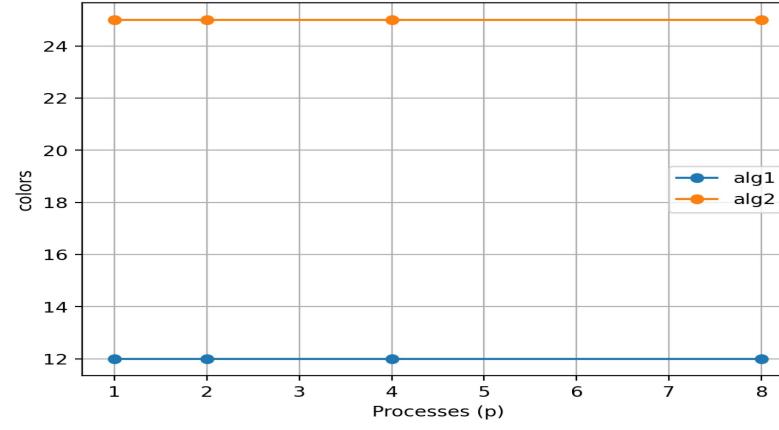
Algorithm Comparison

Algorithm 1 achieves better speedup than Algorithm 2

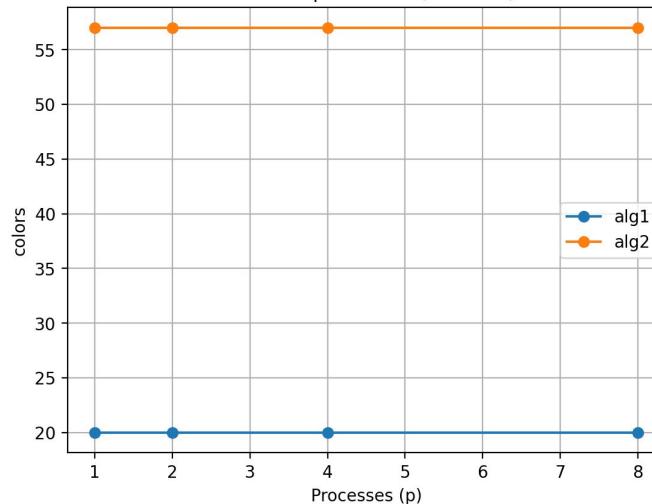
colors vs processes (n=1000)



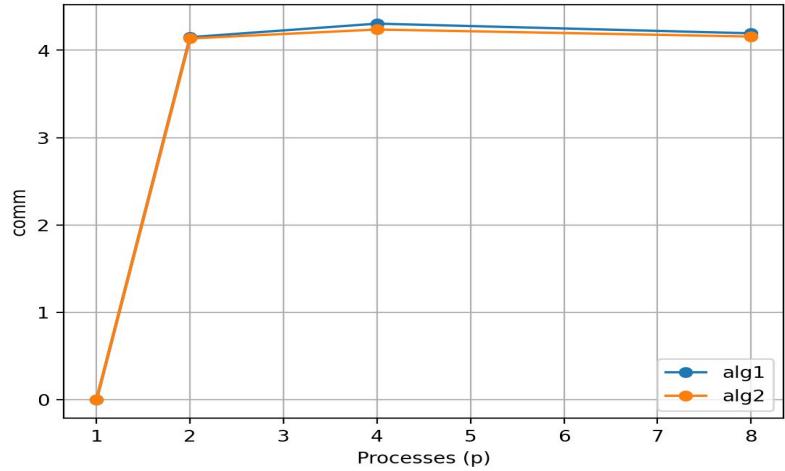
colors vs processes (n=2000)



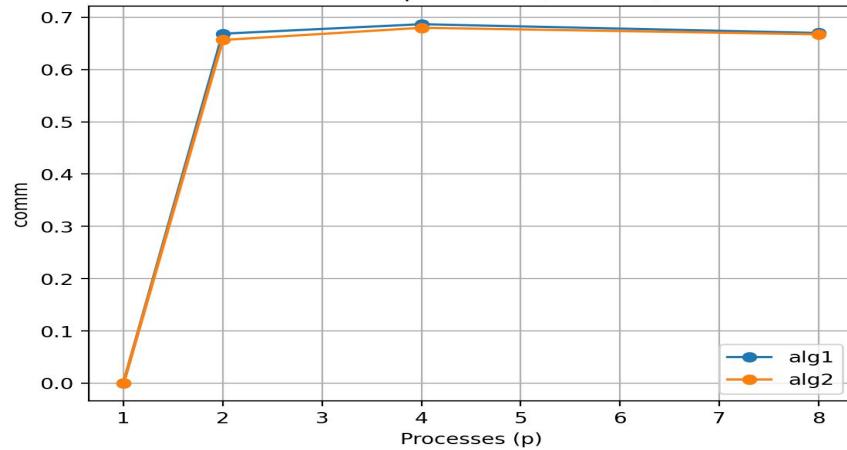
colors vs processes (n=5000)



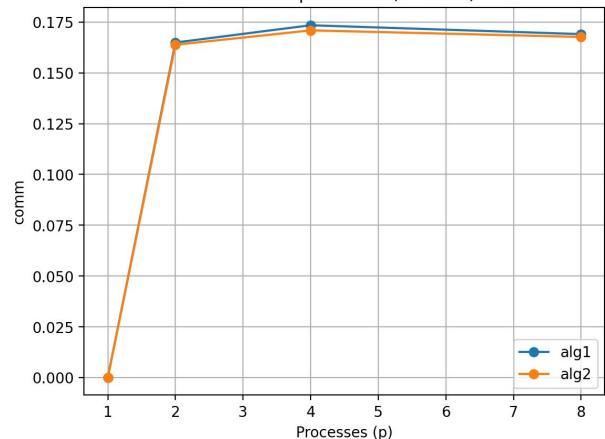
comm vs processes (n=5000)



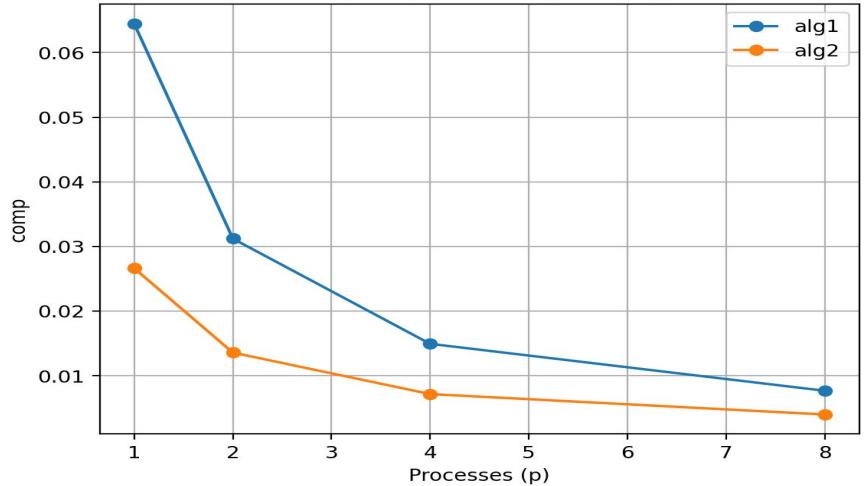
comm vs processes (n=2000)



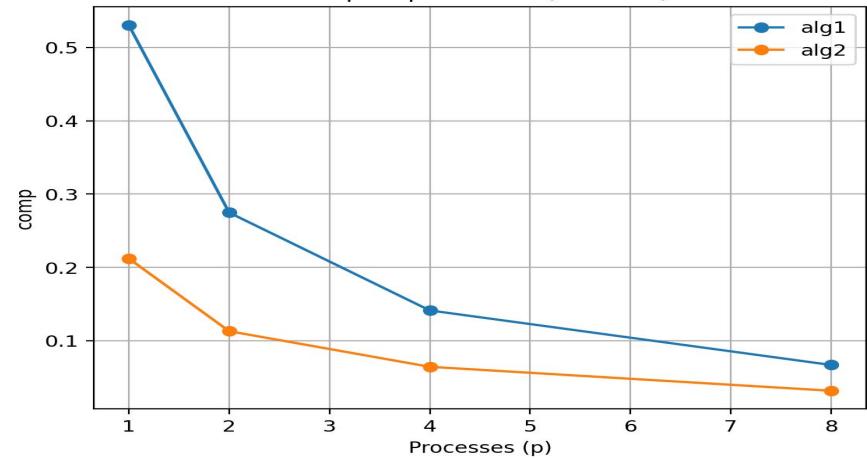
comm vs processes (n=1000)



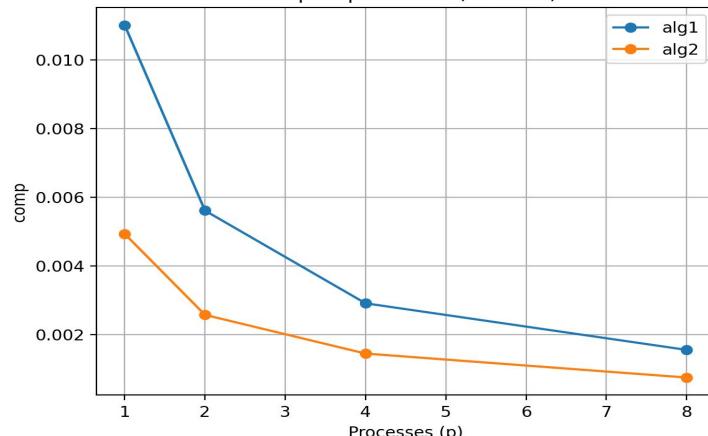
comp vs processes (n=2000)



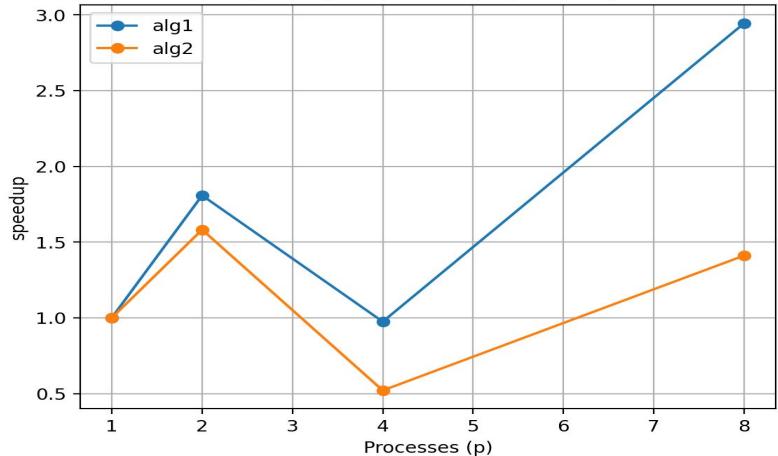
comp vs processes (n=5000)



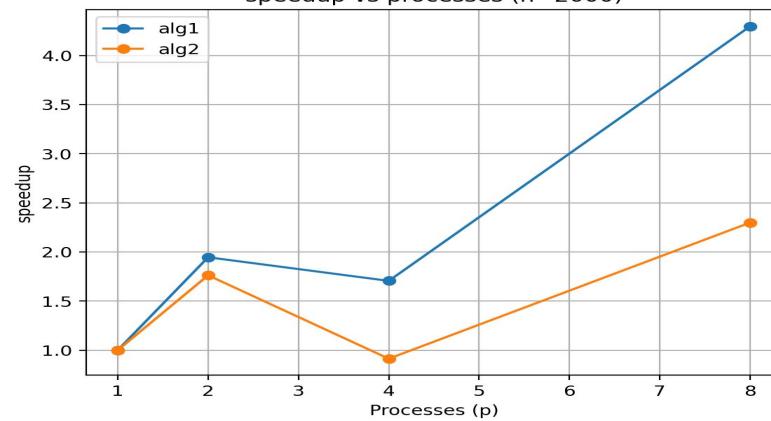
comp vs processes (n=1000)



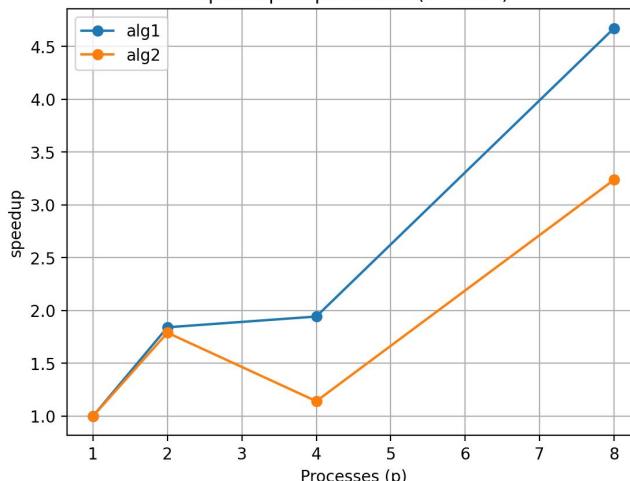
speedup vs processes (n=1000)



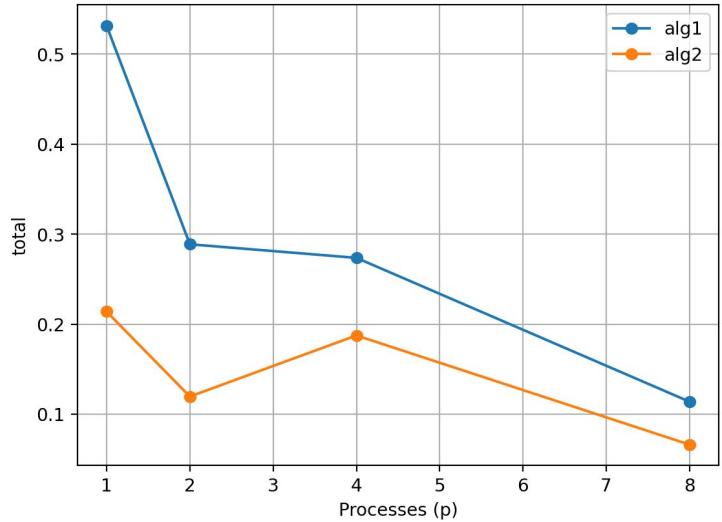
speedup vs processes (n=2000)



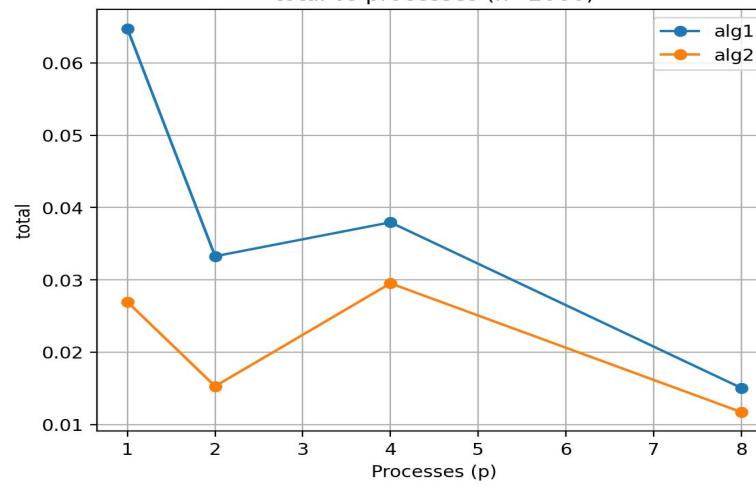
speedup vs processes (n=5000)



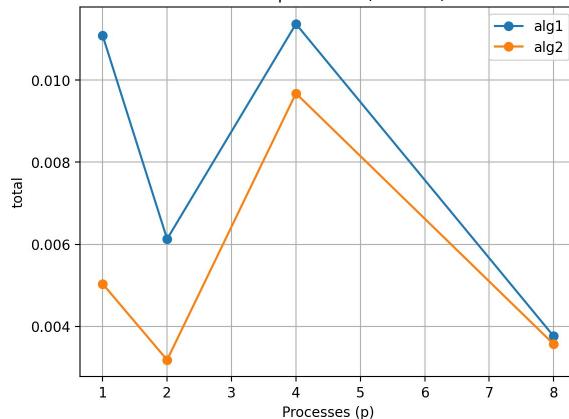
total vs processes (n=5000)



total vs processes (n=2000)



total vs processes (n=1000)



Algorithm Comparison



Algorithm 1: Partitioning Strategy

+ Pros

- ✓ Lower number of rounds
- ✓ Usually faster in practice
- ✓ Often uses fewer colors

- Cons

- ❗ More complex local decision
- ❗ Check neighbors' colors + priorities



Algorithm 2: Divide-and-Conquer Strategy

+ Pros

- ✓ Simple concept: MIS → color → remove
- ✓ Naturally fits "divide-and-conquer" idea

- Cons

- ❗ Typically more rounds (one MIS per color)
- ❗ More colors used than greedy partitioning
- ❗ Slightly heavier communication

Conclusions & Future Work



Conclusions

- _both algorithms produce valid graph colorings in parallel
- Partitioning Strategy: Better execution time, Good speedup with MPI
- Divide-and-Conquer: Conceptually clean MIS-based approach, Slower but easier to reason about



Limitations

- Use of adjacency matrix limits very large graphs
- MPI communication overhead at high process counts



Future Work

- Use adjacency lists / CSR format for sparse graphs
- Better partitioning / load balancing
- Hybrid MPI + OpenMP model
- Test on real-world graph data

Thank you for your attention
Questions?